



Wojciech Krzywiec

Follow

Apr 5, 2020 · 14 min read · Listen

Save



Deployment of multiple apps on Kubernetes cluster — Walkthrough

With this blog post I would like to show you how you can deploy couple applications on minikube (local Kubernetes) cluster.



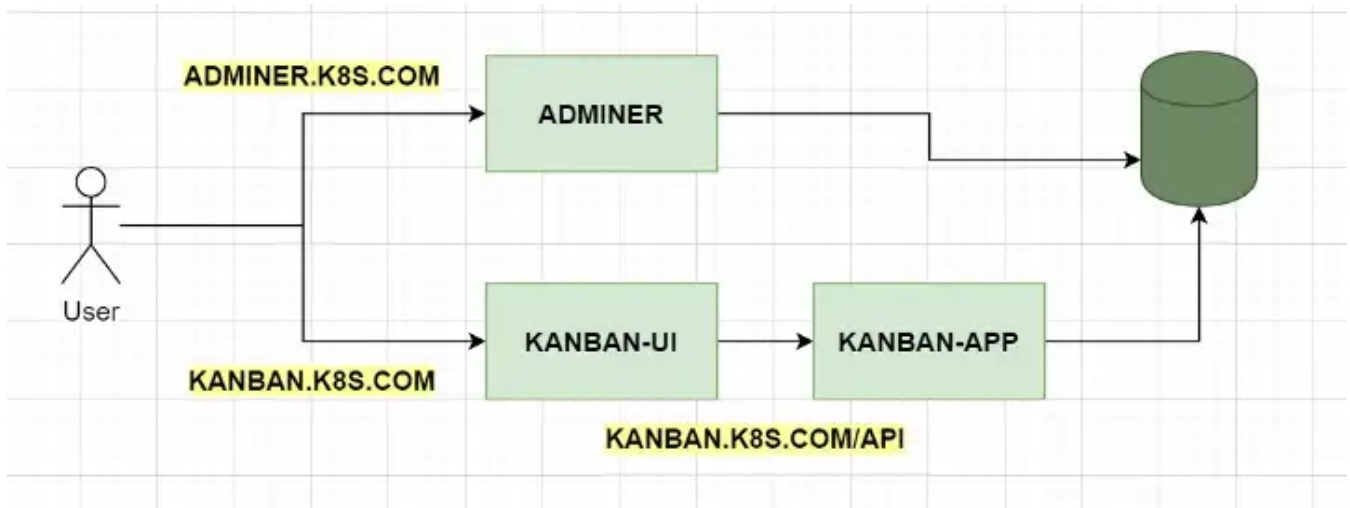
Photo by [Joseph Barrientos](#) on [Unsplash](#)

This is part one of my new series where I compare how to run applications on Kubernetes cluster using 3 approaches:

- kubectl (this one),
- Helm — [*How to deploy application on Kubernetes with Helm*](#),
- Helmfile — [*How to declaratively run Helm charts using helmfile*](#).

Architecture

Before making hands dirty let's see the overall architecture that we want to deploy:



It's based on my previous project — [kanban-board](#), and include 3 services:

- database,
- backend service (*kanban-app*, written in Java with Spring Boot)
- and frontend (*kanban-ui*, written with Angular framework).

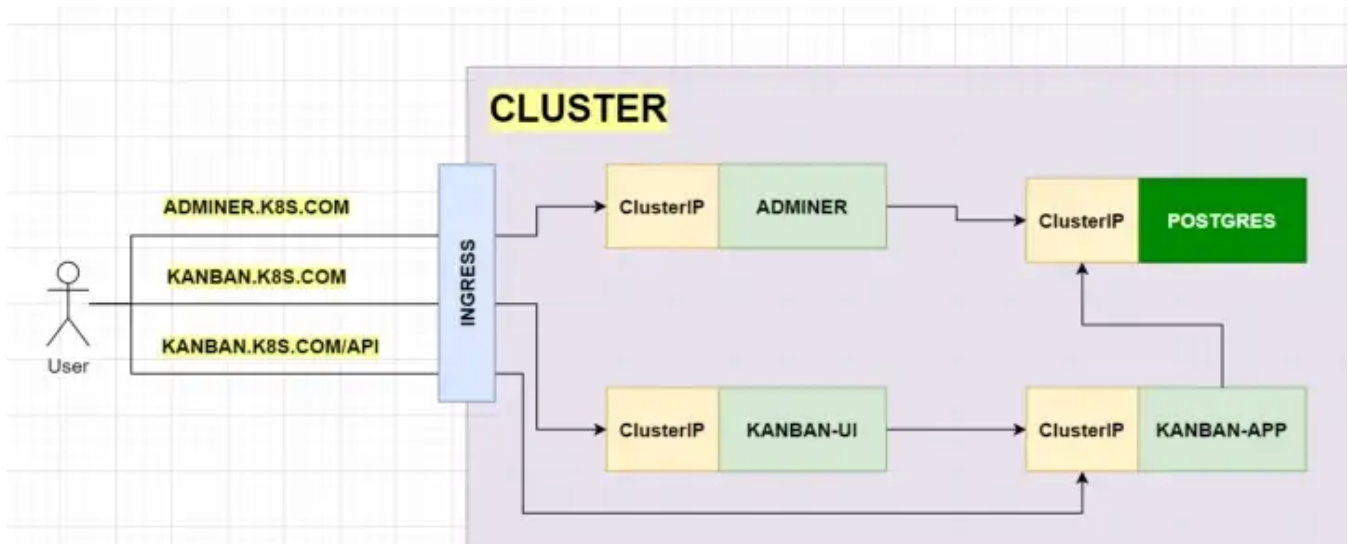
Apart from them I want to deploy the Adminer — UI application to get inside a database.

To enter one of these two UI apps user will need to type one of following URLs in the web browser:

- *kanban.k8s.com*
- *adminer.k8s.com*

The picture above is simplified, just for you to understand the main idea behind this project. Unfortunately it doesn't contain any information of what kind of

Kubernetes Objects we need to create.



If you don't know what some of these objects are, like *Ingress* or *ClusterIP*, don't worry. I'll explain all of that in a minute 😊.

Install Docker, kubectl & minikube

First you need to install all necessary dependencies. Here are links to official documentations which are covering most of popular OSes:

- [Docker](#) (container daemon),
- [kubectl](#) (a CLI tool to interact with cluster),
- [minikube](#) (locally installed, lightweight *Kubernetes* cluster).



Start minikube

Once you've got everything installed you can start the *minikube* cluster by running the CLI command in terminal:

```
$ minikube start
```

```

😊 minikube v1.25.2 on Ubuntu 20.04 (amd64)
✨ Automatically selected the docker driver
👍 Starting control plane node minikube in cluster minikube
🚚 Pulling base image ...
🔥 Creating docker container (CPUs=2, Memory=2200MB) ...
🌊 Preparing Kubernetes v1.23.3 on Docker 20.10.12 ...
  ▪ kubelet.housekeeping-interval=5m
  ▪ Generating certificates and keys ...
  ▪ Booting up control plane ...
  ▪ Configuring RBAC rules ...
🔍 Verifying Kubernetes components...
```

- Using image `gcr.io/k8s-minikube/storage-provisioner:v5`
-  Enabled addons: default-storageclass, storage-provisioner
 Done! `kubectl` is now configured to use "minikube" cluster and "default" namespace by default

To check the status of the cluster:

```
$ minikube status
host: Running
kubelet: Running
apiserver: Running
kubeconfig: Configured
```

To check that `kubectl` is properly configured:

```
$ kubectl cluster-info
Kubernetes master is running at https://127.0.0.1:32768
KubeDNS is running at
https://127.0.0.1:32768/api/v1/namespaces/kube-system/services/kube-dns:dns/proxy

To further debug and diagnose cluster problems, use 'kubectl
cluster-info dump'.
```

Modify hosts file

To make the <http://adminer.k8s.com> & <http://kanban.k8s.com> work you need to edit the **hosts** file on your PC.

The location of it depends on the OS:

- [Linux \(Ubuntu\)](#)
- [Windows 10](#)
- [Mac](#)

When you find it, add following lines:

```
127.0.0.1    adminer.k8s.com
127.0.0.1    kanban.k8s.com
```

It will map your `localhost` IP address to both hostnames and makes them accessible after running the `minikube tunnel` command.

Add Adminer

Finally everything is set up and we can start with deploying applications. First one will be *Adminer* app.

In *Kubernetes* world the smallest deployable object is a **Pod**. It can hold one or more (Docker, cri-o) containers and also has some metadata information (e.g name, labels) that are needed. Sometimes *Pods* can be treated as single applications, because they usually have only one single container inside.

But we won't create *Pods* in this exercise 🤪. Not directly at least 😊.

The problem with them is that if you're creating them manually you won't be able to easily scale their number. Also if your application inside the *Pod* crashes your *Pods* also crashes and there is no mechanism to restart it again.

Luckily there is a **Deployment** for help 🍷.

In order to create it for *Adminer* you need to have a file called **adminer-deployment.yaml** which is defined as follows:

```
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: adminer
5    labels:
6      app: adminer
7      group: db
8  spec:
9    replicas: 1
10   selector:
11     matchLabels:
12       app: adminer    # indicates which Pods (with which labels) need be handled by
this Deployment
13   template:
14     metadata:          # indicates the metadata that will be added to each Pod
15       labels:
16         app: adminer
17         group: db
18     spec:
19       containers:      # defines a running container inside the Pod
20         - name: adminer
21           image: adminer:4.7.6-standalone
22           ports:
23             - containerPort: 8080    # which port of the container is exposed to the
Pod
24           env:
25             - name: ADMINER_DESIGN
26               value: pepa-linha
27             - name: ADMINER_DEFAULT_SERVER
28               value: postgres
29           resources:
30             limits:
31               memory: "256Mi"
32               cpu: "500m"
```

adminer-deployment.yaml hosted with ❤ by GitHub

[view raw](#)

First section is responsible for defining of what kind of object we're creating (`apiVersion` , `kind`) followed by some metadata including name & labels (`metadata`).

Next section — `spec` — is called specification where we define specifications of a *Deployment*:

- `replicas` — indicates how many *Pods* of the same type will be created,

- `selector.matchLabels` — defines how *Deployment* will find *Pods* that it needs to take care of, in this case it will look for a Pod which is labeled with `app: adminer` ,
- `template.metadata` — tells what metadata will be added to each *Pod*, in this case all of them will have `labels : app: adminer , group: db .`
- `template.spec.containers` — is a list of containers that will be inside a *Pod*. In this case I put only one container, which is based on `adminer:4.7.6-standalone` Docker image and exposes `containerPort: 8080` . Moreover with `env` section we inject environment variable to the container to configure *Adminer UI* (full documentation can be found [here](#)). And finally we decide how much RAM and CPU an will require.

Now you can run following command in the terminal:

```
$ kubectl apply -f adminer-deployment.yaml
deployment.apps/adminer created
```

To check if everything is ok you can run:

```
$ kubectl get deployments
NAME          READY    UP-TO-DATE    AVAILABLE    AGE
adminer       1/1      1              1             30s

$ kubectl describe deployment adminer
... many details about the Deployment ...

$ kubectl get pods
NAME                                READY    STATUS    RESTARTS    AGE
adminer-994865d4b-kqck5            1/1      Running    0             24m

$ kubectl describe pod adminer-994865d4b-kqck5
... many details about the Pod ...
```

Great! It worked! But there is a problem. How to open the Adminer page?

To handle this problem we need to use another type of *Kubernetes* object — **Service**.

Per design *Kubernetes* is assigning the IP for each Pod, which might be problematic, because *Pods* don't live forever. Actually they are constantly created and deleted, all

the time. And for each new *Pod* new IP is assigned. And that's creates some kind of networking hell, because other applications inside the cluster would need to update the IP addresses of connected with *Pods* every time new instance is created.

Luckily *Services* are to the rescue. They solve that problem by having a single DNS name for all *Pods* handled by the *Deployment*. So no matter what IP address *Pod* have, all applications are pointing to the *Service* which do all the job - finding the right *Pod*. Plus *Services* are taking care of load balancing of the traffic if there are more than *Pod* replicas.

To create such object add new YAML file with Service definition **adminer-svc.yaml**:

```
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: adminer
5    labels:
6      group: db
7  spec:
8    type: ClusterIP
9    selector:           # indicates into which pods ClusterIP provides access
10     app: adminer
11    ports:
12     - port: 8080       # port exposed outside ClusterIP
13       targetPort: 8080 # port exposed by Pod
```

adminer-svc.yaml hosted with ❤ by GitHub

[view raw](#)

This one is a little bit shorter. But like the last time there is section defining the type of the object and it's metadata. Then there is a `spec` section where couple of properties are set:

- `type: ClusterIP` — indicates what type of the Service we want to deploy. There are several types, but I've decided to use **ClusterIP**. And the main reason for that is because I didn't want to expose every *Pod* outside the cluster. What *ClusterIP* does is that it exposes assigned *Pods* to other *Pods* inside the cluster, but not outside.
- `selector` — here we say to which *Pods* this *Service* provide access, in this case it provide access to a *Pod* with `app: adminer` label.

- `ports` — indicates the mappings of the port exposed by the *Pod* to the *ClusterIP* port which will be available for other applications inside cluster.

And now we can create this Service with command:

```
$ kubectl apply -f adminer-svc.yaml
service/adminer created
```

And to check if everything is working:

```
$ kubectl get svc
NAME         TYPE          CLUSTER-IP    EXTERNAL-IP  PORT(S)    AGE
adminer      ClusterIP     10.99.85.149   <none>       8080/TCP   9s
kubernetes   ClusterIP     10.96.0.1      <none>       443/TCP    3m34s

$ kubectl describe svc adminer
... many details about the ClusterIP...
```

Okay! So am I able now to open the Adminer page?

Of course not 😊. We need to do one more thing.

Add Ingress Controller

As it was mentioned before, *ClusterIP* exposes the app only for other apps inside the cluster. And in order to get to it from outside of it we need to use a different approach.

Here comes Ingress to the rescue, which is a gateway to our cluster. And the object that we need to create is called Ingress Controller and it's an implementation of *Ingress*.

But here is the tricky part. There are lots of *Ingress Controllers* available. Some of them are opensource, but some of them are paid one. For this project I've chosen an "official" Kubernetes Ingress Controller based on Nginx. But please be not confused with another one, also based on Nginx but created by NGINX Inc - this one is paid. Apart from these both there are also other *Ingress Controllers* available like Kong Ingress, or Traefik.

Luckily for us minikube comes with already built-in *Ingress Controller*. The only thing to do is to run following command:

```
$ minikube addons enable ingress
```

🌟 The 'ingress' addon is enabled

To make it work run following command in a separate terminal window:

```
$ minikube tunnel
```

✅ Tunnel successfully started

📌 NOTE: Please *do not* close this terminal as this process must stay alive *for* the tunnel to be accessible ...

So then we can move to defining the routing rule to get inside the *Adminer* web page. Therefore we need to create an *Ingress* object defined in a file **ingress-controller.yaml**:

```
1  apiVersion: networking.k8s.io/v1
2  kind: Ingress
3  metadata:
4    name: ingress-service
5    annotations:
6      kubernetes.io/ingress.class: nginx
7  spec:
8    rules:
9      - host: adminer.k8s.com
10      http:
11        paths:
12          - path: /
13            pathType: Prefix
14            backend:
15              service:
16                name: adminer
17                port:
18                  number: 8080
```

ingress-controller.yaml hosted with ❤ by GitHub

[view raw](#)

As usual, first there is a definition of kind of the *Kubernetes* object we want to create. Then it's followed by `metadata` with the name of the object as usual and also add a

new section — annotations .

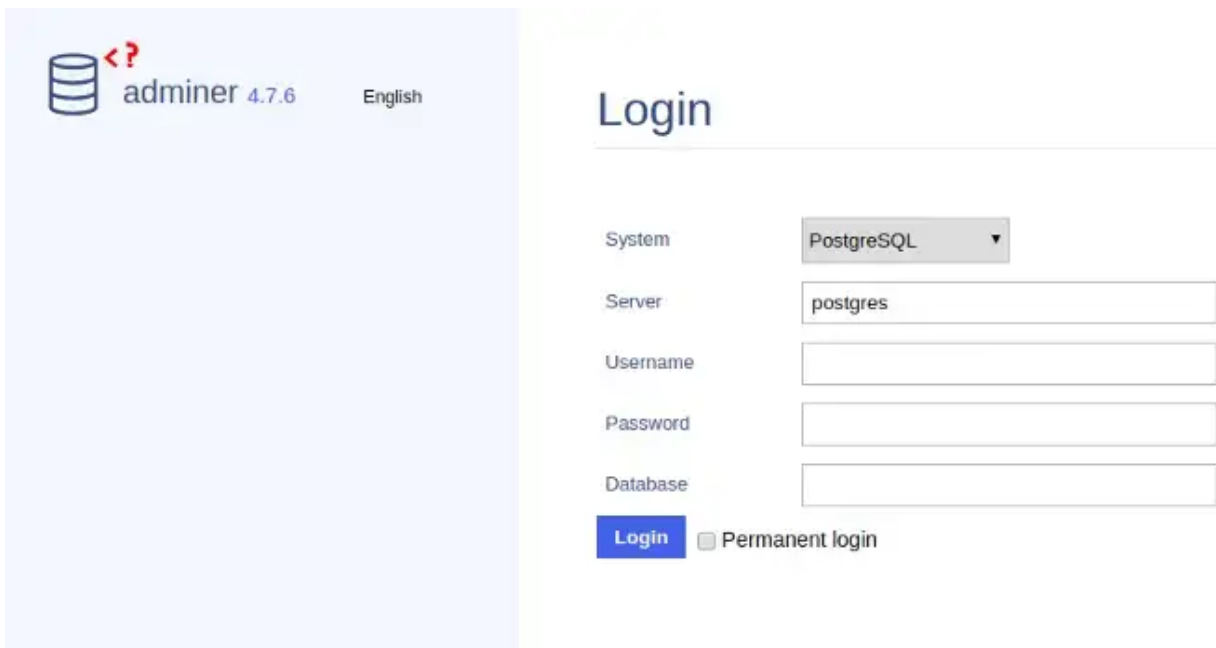
This one is very important for *Ingress* because with that we can configure its behavior. In my example, I've used the most simple one, but there are a lot more of possibilities.

And finally there is a `spec` section where we provided first rule, that all requests from the host `adminer.k8s.com` **will be routed to the ClusterIP with a name `adminer`** .
ClusterIP, neither *Deployment* nor *Pod*!

After applying it into the cluster:

```
$ kubectl apply -f ingress-controller.yaml
ingress.networking.k8s.io/ingress-controller created
```

And finally after typing <http://adminer.k8s.com> in a web browser this page show up:



Awesome! But how to login to the database? Wait, but what database? We don't have any database at all!

Add PostgreSQL database

Right, we need to set up our database. To do that we need to create another pair of *Deployment-ClusterIP*, but this time with *PostgreSQL*.

And here, again is a tricky part. Databases are not the “usual” stateless services, they store information and should not be as easily killed as “regular” *Pods* are. Even if it crashes we want to have data persisted somewhere. Therefore we need to create a space (directory) on our disk, which will be accessible by *PostgreSQL* container and will be outside the *Kubernetes* cluster.

In order to do that we need to create a new type of object called — PersistentVolumeClaim. It provides some storage located on our computer (or server) for *Pods*.

To create it once again we create a YAML file with a name — **postgres-pvc.yaml**:

```
1  apiVersion: v1
2  kind: PersistentVolumeClaim
3  metadata:
4    name: postgres-persistent-volume-claim
5  spec:
6    accessModes:
7      - ReadWriteOnce
8    resources:
9      requests:
10       storage: 4Gi
```

postgres-pvc.yaml hosted with ♥ by GitHub

[view raw](#)

Again first sections include the definition of the type of object we want to create together with some metadata. Then in the `spec` section we tell *Kubernetes* that this *Volume* has read-write access right and we want to use 4GiB of memory (I know maybe that's too much for such small example).

And after applying it in the terminal:

```
$ kubectl apply -f postgres-pvc.yaml
persistentvolumeclaim/postgres-persistent-volume-claim created
```

And to find out if everything is ok:

```
$ kubectl get pvc
NAME          STATUS  VOLUME  CAPACITY  ACCESS MODE  STORAGECLASS  AGE
postgres...  Bound  pvc-43  4Gi       RWO          standard      40s
```

```
$ kubectl describe pvc postgres-persistent-volume-claim
... many details about the PersistentVolumeClaim...
```

Next, we should be able to create *Deployment & ClusterIP* for *PostgreSQL*, but first I would like to introduce new type of *Kubernetes* object, which will hold some configuration values and is called — ConfigMap.

This type of object is very useful when we want to inject environment variables to multiple containers in the *Pods*. It makes configuration of multiple *Pods/Deployments* very clean because we can have a single point of truth for our configuration. And if we decide to change it, we can do that in one place.

In this project I want keep the database configuration in *ConfigMap*, because I want to pass database config values to two *Pods* — one for postgres *Deployment*, and one for the backend service.

The definition of the *ConfigMap* is in the **postgres-config.yaml** file and is as follows:

```
1  apiVersion: v1
2  kind: ConfigMap
3  metadata:
4    name: postgres-config
5    labels:
6      group: db
7  data:
8    POSTGRES_DB: kanban
9    POSTGRES_USER: kanban
10   POSTGRES_PASSWORD: kanban
```

postgres-config.yaml hosted with ♥ by GitHub

[view raw](#)

Except the usual sections — `apiVersion` , `kind` and `metadata` there is a new one instead of `spec` — `data` . It's where there are pairs of keys & values for environment variables that we will be injecting to the containers.

To create this object we need to run the command:

```
$ kubectl apply -f postgres-config.yaml
configmap/postgres-config created
```

And to check it we can run commands:

```
$ kubectl get configmap
```

NAME	DATA	AGE
postgres-config	3	2m31s

```
$ kubectl describe configmap postgres-config
```

```
... many details about the ConfigMap...
```

Now we can move on to the definition of PostgreSQL *Deployment* — **postgres-deployment.yaml**:

```

1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: postgres
5    labels:
6      app: postgres
7      group: db
8  spec:
9    replicas: 1
10   selector:
11     matchLabels:
12       app: postgres
13   template:
14     metadata:
15       labels:
16         app: postgres
17         type: db
18     spec:
19       volumes:                                # indicates which PVC are available for this
Deployment
20         - name: postgres-storage
21           persistentVolumeClaim:
22             claimName: postgres-persistent-volume-claim
23       containers:
24         - name: postgres
25           image: postgres:9.6-alpine
26           ports:
27             - containerPort: 5432
28           envFrom:
29             - configMapRef:
30               name: postgres-config
31           volumeMounts:                        # indicates which Volume (from
spec.template.spec.volumes) should be used
32             - name: postgres-storage          # name of the Volume
33               mountPath: /var/lib/postgresql/data # path inside the container
34       resources:
35         limits:
36           memory: "256Mi"
37           cpu: "500m"

```

postgres-deployment.yaml hosted with ❤ by GitHub

[view raw](#)

As most of the parts were already discussed I'll skip them and move to new ones:

- `spec.template.spec.volumes` — here we're adding created PVC to the *Deployment*, so all containers inside of it will be able to use it,

- `spec.template.spec.containers[0].image` — here we specify what Docker image we want to use for our database,
- `spec.template.spec.containers[0].envFrom` — indicates from which *ConfigMap* we want to inject environment variables,
- `spec.template.spec.containers[0].volumeMounts` — tells *Kubernetes* which *Volume* to use (defined in the `spec.template.spec.volumes` section) and map it to a particular folder inside the container — basically all data inside the `mountPath` will be stored outside the cluster.

Similarly, we define the *ClusterIP* with a file `postgres-svc.yaml`:

```
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: postgres
5    labels:
6      group: db
7  spec:
8    type: ClusterIP
9    selector:
10     app: postgres
11    ports:
12     - port: 5432
13       targetPort: 5432
```

`postgres-svc.yaml` hosted with ♥ by GitHub

[view raw](#)

There is nothing new here, except the port mapping which is specific for PostgreSQL.

To create both objects we can run:

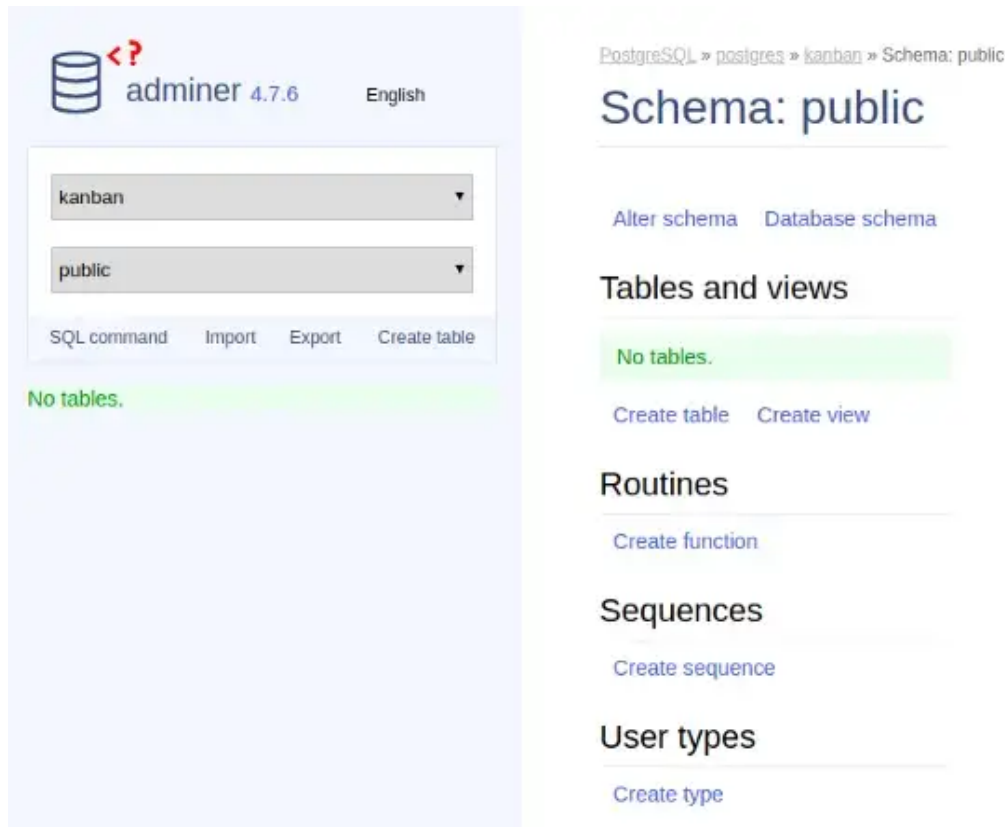
```
$ kubectl apply -f postgres-deployment.yaml
deployment.apps/postgres created

$ kubectl apply -f postgres-svc.yaml
service/postgres created
```

And now if you go to the *Adminer* once again, type following credentials:

System: PostgreSQL
Server: postgres
Username: kanban
Password: kanban
Database: kanban

You should be able to login to a page:



Awesome! The database is set up, so we can move on to kanban-app (backend) and kanban-ui (frontend) services.

Add kanban-app

First let's provide all necessary definitions for backend service. As it was for *Adminer*, we need also to have create *Deployment* and *Service* for backend service.

Therefore, the **kanban-app-deployment.yaml** file looks as follows:

```
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: kanban-app
5    labels:
6      app: kanban-app
7      group: backend
8  spec:
9    replicas: 1
10   selector:
11     matchLabels:
12       app: kanban-app
13   template:
14     metadata:
15       labels:
16         app: kanban-app
17         group: backend
18     spec:
19       containers:
20         - name: kanban-app
21           image: wkrzywiec/kanban-app:k8s
22           ports:
23             - containerPort: 8080
24           envFrom:
25             - configMapRef:
26                 name: postgres-config
27           env:
28             - name: DB_SERVER
29               value: postgres
30           resources:
31             limits:
32               memory: "256Mi"
33               cpu: "500m"
```

kanban-app-deployment.yaml hosted with ❤ by GitHub

[view raw](#)

In the container specification I provided my own Docker image which I've published on Docker Hub. It exposes port 8080 and uses some of the environment variables located either in *ConfigMap* (`envFrom.configMapRef`) or from manually added environment variable only for this *Deployment*— `env` .

Next, we define the **kanban-app-svc.yaml** file:

```
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: kanban-app
5    labels:
6      group: backend
7  spec:
8    type: ClusterIP
9    selector:
10     app: kanban-app
11    ports:
12     - port: 8080
13       targetPort: 8080
```

kanban-app-svc.yaml hosted with ♥ by GitHub

[view raw](#)

There are no new things there in compare to previous *Services*.

To apply both definitions we need to run the commands:

```
$ kubectl apply -f kanban-app-deployment.yaml
deployment.apps/kanban-app created

$ kubectl apply -f kanban-app-svc.yaml
service/kanban-app created
```

Now you would want to test it, but in order to do so we need to configure the Ingress Controller so we can enter the [Swagger UI](#) page to check the API of the backend service.

We need to add a new host to the `ingress-controller.yaml` file so it will look as follows:

```
1  apiVersion: networking.k8s.io/v1
2  kind: Ingress
3  metadata:
4    name: ingress-service
5    annotations:
6      kubernetes.io/ingress.class: nginx
7  spec:
8    rules:
9      - host: adminer.k8s.com
10      http:
11        paths:
12          - path: /
13            pathType: Prefix
14            backend:
15              service:
16                name: adminer
17                port:
18                  number: 8080
19      - host: kanban.k8s.com
20      http:
21        paths:
22          - path: /api/
23            pathType: Prefix
24            backend:
25              service:
26                name: kanban-app
27                port:
28                  number: 8080
```

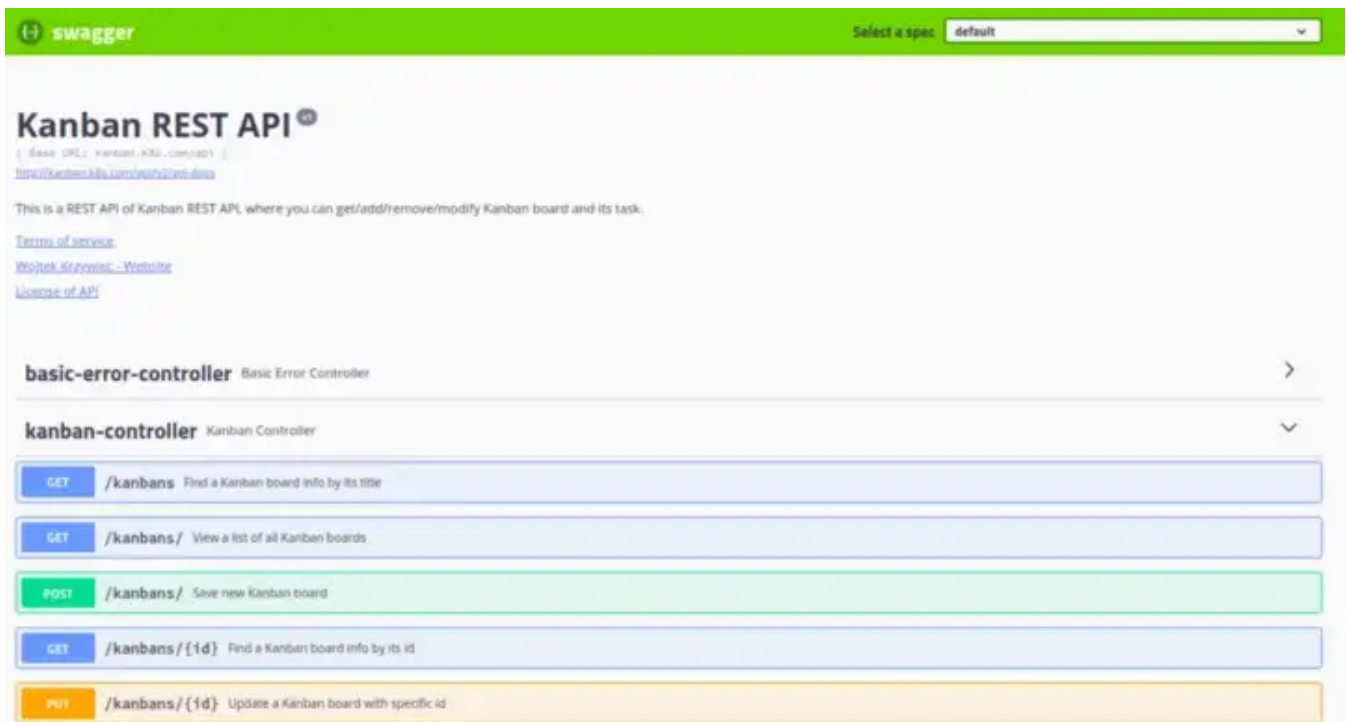
ingress-controller.yaml hosted with ❤ by GitHub

[view raw](#)

And then apply those changes to the cluster by running the command:

```
$ kubectl apply -f ingress-controller.yaml
ingress.networking.k8s.io/ingress-service configured
```

Now, if you enter the <http://kanban.k8s.com/api/swagger-ui.html> address in the web browser you should get the overview of the REST API that this application is providing.



You can also go to the Adminer (<http://adminer.k8s.com>) and check if new tables were added to the database (they were added by the Liquibase script during start up of kanban-app).

PostgreSQL » postgres » kanban » Schema: public

Logout

Schema: public

[Alter schema](#) [Database schema](#)

Tables and views

Search data in tables (4)

	Table	Engine	Collation	Data Length [?]	Index Length [?]	Data Free	Auto Increment	Rows [?]	Comment [?]
<input type="checkbox"/>	databasechangelog	table		8,192	8,192	?	?	0	
<input type="checkbox"/>	databasechangeloglock	table		8,192	16,384	?	?	0	
<input type="checkbox"/>	kanban	table			16,384	?	?	0	
<input type="checkbox"/>	task	table			16,384	?	?	0	

Selected (0)

Move to other database:

Add kanban-ui

And at last, we can add the UI application. Again, we need to define the *Deployment* and *ClusterIP*.

Here is the **kanban-ui-deployment.yaml** file

```
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: kanban-ui
5    labels:
6      app: kanban-ui
7      group: frontend
8  spec:
9    replicas: 1
10   selector:
11     matchLabels:
12       app: kanban-ui
13   template:
14     metadata:
15       labels:
16         app: kanban-ui
17         group: frontend
18     spec:
19       containers:
20         - name: kanban-ui
21           image: wkrzywiec/kanban-ui:k8s
22           ports:
23             - containerPort: 80
24           resources:
25             limits:
26               memory: "256Mi"
27               cpu: "500m"
```

kanban-ui-deployment.yaml hosted with ❤ by GitHub

[view raw](#)

And **kanban-ui-svc.yaml** file:


```
1  apiVersion: v1
2  kind: Service
3  metadata:
4    name: kanban-ui
5    labels:
6      group: backend
7  spec:
8    type: ClusterIP
9    selector:
10     app: kanban-ui
11    ports:
12     - port: 80
13       targetPort: 80
```

kanban-ui-svc.yaml hosted with ♥ by GitHub

[view raw](#)

Nothing special in both files, so we can go right away to applying both of them to the cluster:

```
$ kubectl apply -f kanban-ui-deployment.yaml
deployment.apps/kanban-ui created
```

```
$ kubectl apply -f kanban-ui-svc.yaml
service/kanban-ui created
```

And again, to test it we need to expose it outside cluster. For that we need to configure *Ingress Controller*:

```
1  apiVersion: networking.k8s.io/v1
2  kind: Ingress
3  metadata:
4    name: ingress-service
5    annotations:
6      kubernetes.io/ingress.class: nginx
7  spec:
8    rules:
9      - host: adminer.k8s.com
10      http:
11        paths:
12          - path: /
13            pathType: Prefix
14            backend:
15              service:
16                name: adminer
17                port:
18                  number: 8080
19      - host: kanban.k8s.com
20      http:
21        paths:
22          - path: /api/
23            pathType: Prefix
24            backend:
25              service:
26                name: kanban-app
27                port:
28                  number: 8080
29          - path: /
30            pathType: Prefix
31            backend:
32              service:
33                name: kanban-ui
34                port:
35                  number: 80
```

ingress-controller.yaml hosted with ❤ by GitHub

[view raw](#)

And now, if you open the address — <http://kanban.k8s.com> you should get this page:



You can now add Kanban boards, tasks, etc.

But one more point before the wrap up.

How the kanban-ui is connected with kanban-app?

The answer to this question is in the configuration file of Nginx server, included in the Docker image of *kanban-ui* — [default.conf](#).

```
1  server {
2      listen 80;
3      server_name kanban-ui;
4      root /usr/share/nginx/html;
5      index index.html index.html;
6
7      location /api/kanbans {
8          proxy_pass http://kanban-app:8080/api/kanbans;
9      }
10
11     location /api/tasks {
12         proxy_pass http://kanban-app:8080/api/tasks;
13     }
14
15     location / {
16         try_files $uri $uri/ /index.html;
17     }
18 }
```

default.conf hosted with ♥ by GitHub

[view raw](#)

In above example the address <http://kanban-app:8080> is a DNS address of the *ClusterIP*, not the *Deployment*.

Conclusion

With this blog post I've tried to walk you through all the steps to deploy couple applications into a local *Kubernetes* cluster.

But there is one problem. How to avoid creating such great number of YAML files? And is it a single command with which we could deploy all these objects all at once?

For a second question there is a simple answer — you can run the `kubectl apply` command not on every single file but on the entire folder where they are located i.e.:

```
$ kubectl apply -f ./k8s
deployment.apps/adminer created
service/adminer created
ingress.networking.k8s.io/ingress-service created
deployment.apps/kanban-app created
service/kanban-app created
deployment.apps/kanban-ui created
service/kanban-ui created
configmap/postgres-config created
deployment.apps/postgres created
persistentvolumeclaim/postgres-persistent-volume-claim created
service/postgres created
```

But for the first question, how to avoid such boilerplate code there is no simple question. But I'll try to address it in my next post, where I'll deploy same services,

[Open in app](#) ↗

[Get unlimited access](#)



As usual here are links to my repositories, first with all Kubernetes YAML files:

wkrzywiec/k8s-helm-helmfile

Contribute to wkrzywiec/k8s-helm-helmfile development by creating an account on GitHub.

github.com

And second with the source code of kanban-app & kanban-ui:

wkrzywiec/kanban-board

This is a simple implementation of a Kanban Board, a tool that helps visualize and manage work. Originally it was first...

github.com

30th May 2022 update: Couple of things were updated or removed to comply with the latest version of minikube (1.25) and Kubernetes (1.23) including Ingress Controller definition, editing hosts file, adding resource limits to Deployment definitions. Thanks [Angelos](#) and [Arkadiusz Halicki](#) for catching some of them!

References

Kubernetes Documentation

Kubernetes is an open source container orchestration engine for automating deployment, scaling, and management of...

kubernetes.io

Installing Kubernetes with Minikube

 313 |  12 | 

Edit This Page Minikube is a tool that makes it easy to run Kubernetes locally. Minikube runs a single-node Kubernetes...

kubernetes.io

Kubernetes Ingress 101: NodePort, Load Balancers, and Ingress Controllers

This article will introduce the three general strategies in Kubernetes for ingress, and the tradeoffs with each...

blog.getambassador.io

Studying the Kubernetes Ingress system

I have been researching how the Kubernetes Ingress system works. My use case is to setup an autoscaled Nginx cluster...

www.joyfulbikeshedding.com

Overview of kubectl

Edit This Page Kubectl is a command line tool for controlling Kubernetes clusters. kubectl looks for a file named...

kubernetes.io

kubectl apply vs kubectl create?

kubectl run = kubectl create deployment Simple, easy to learn and easy to remember. Require only a single step to make...

stackoverflow.com

Using Kubernetes to Deploy PostgreSQL

Kubernetes is an open source container orchestration system for automating deployment, scaling and management of...

severalnines.com

Dev Ops

Kubernetes

Cloud Computing

Java

Postgres

Stay tune for upcoming publication!

Subscribe to my feed to get notification for new publications

Emails will be sent to hamdi.bouhani@dealroom.co. [Not you?](#)

